


# COMP810 Data Warehousing and Big Data

Map-Reduce and Hadoop


Dr Weihua Li

# Outline

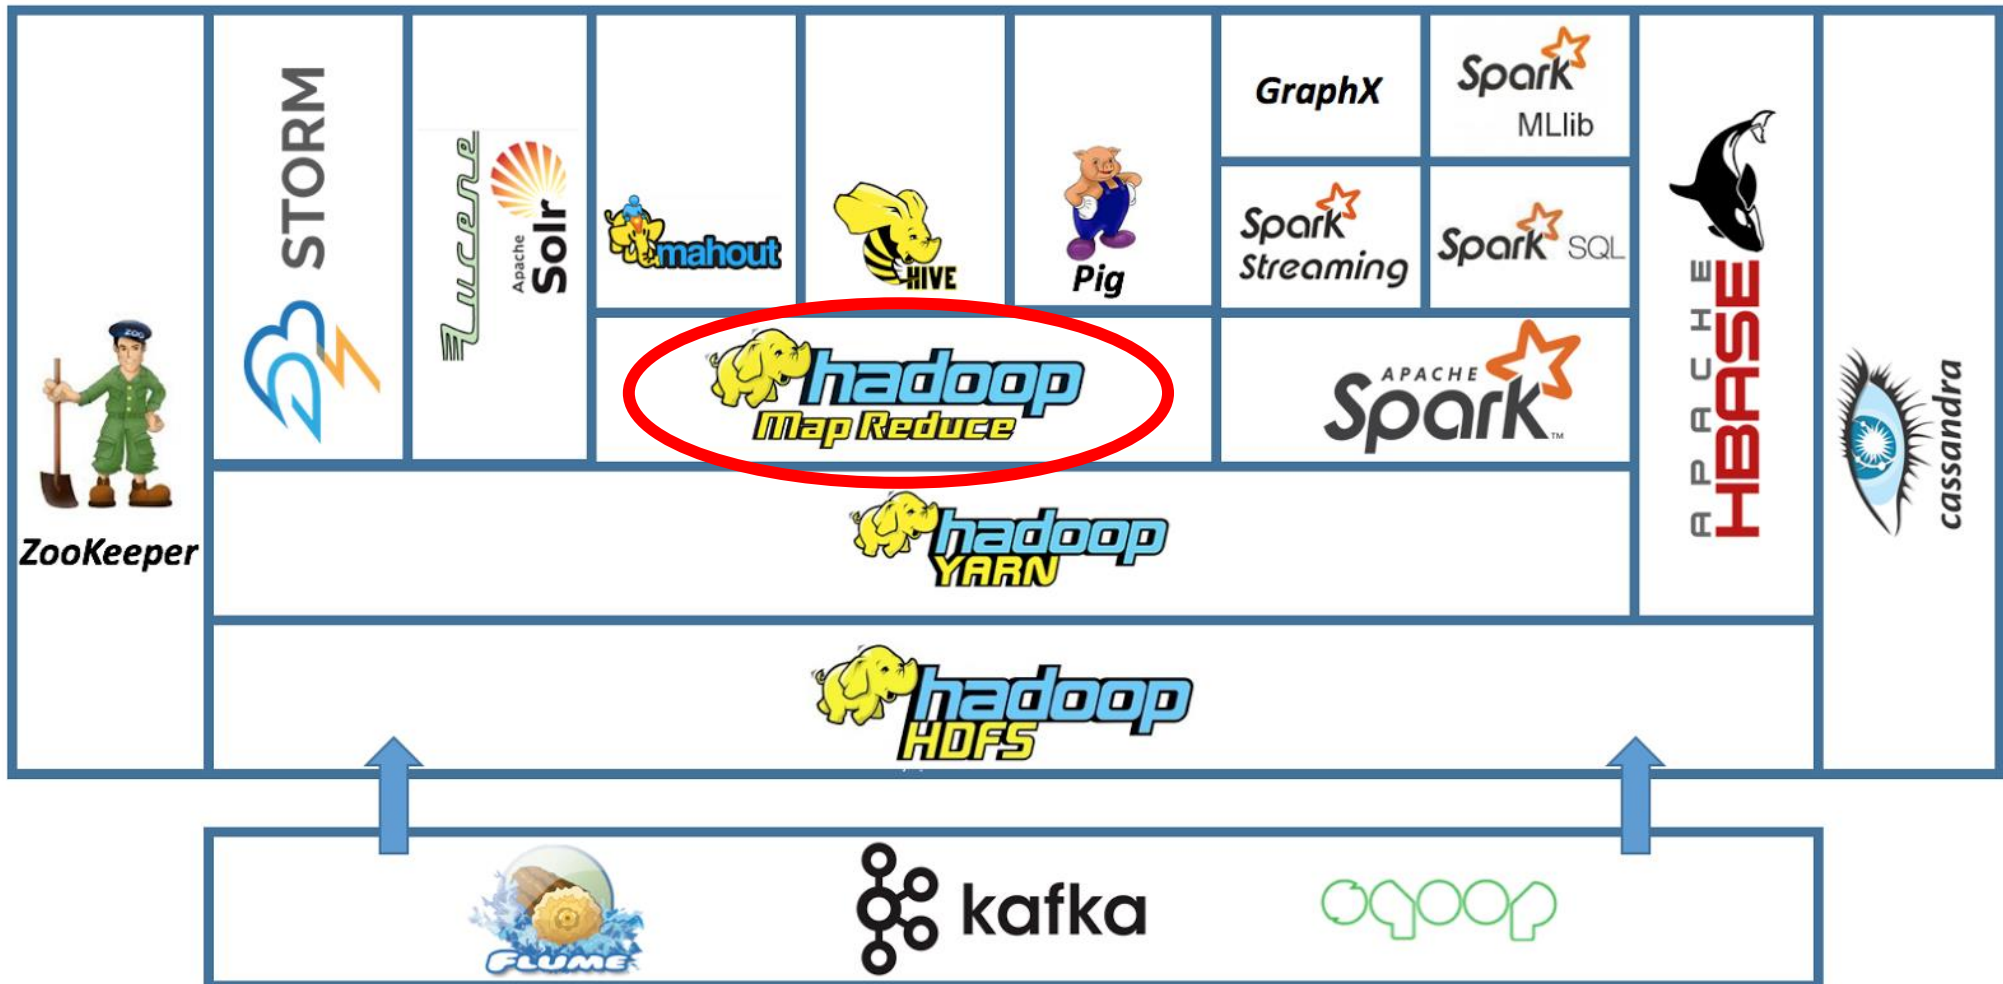
- MapReduce
- Apache Hadoop



# MapReduce



# Apache Hadoop Ecosystem



# Motivation: Google Example

20+ billion web pages x 20 KB = 400+ TB

1 computer reads 3.5 GB/sec from disk

- 32 hours (1.3 days) to read the web



Takes even longer to do something useful with the data!

Today, a standard architecture for such problems is:


- Cluster of commodity Linux nodes
- Commodity network (ethernet) to connect them

# Cluster Architecture



2-10 Gbps backbone between racks

1 Gbps between  
any pair of nodes  
in a rack



Each rack contains 16-64 nodes.

In 2011 it was estimated that Google had 1M machines.



# Cluster Computing Challenges

- Node Failures
  - One server may stay up 3 years (1,000 days)
  - If you have 1,000 servers, expect to have 1 failure/day
  - 1M servers in a cluster, expect to have 1,000 failures every day!
- Network Bottleneck
  - Network bandwidth = 1Gbps
  - Moving 10TB takes 1 day
- Distributed Programming is hard
  - Need a simple model that hides most of the complexity



# MapReduce

- Map-reduce addresses all of these three challenges
  - **Node Failures:** Store data redundantly on multiple nodes for persistence and availability.
  - **Network Bottleneck:** Move computation close to data to minimize data movement.
  - **Uneasy Distributed Programming:** Simple programming model to hide the complexity



# Redundant Storage Infrastructure

- **Distributed File System**


- Store data across a cluster, and each piece of data **multiple times**.
- Provide global file namespace, redundancy and availability.
- E.g., Google File System (GFS), Hadoop File System (**HDFS**)

- **Typical Usage Pattern**

- Huge files (100s of GB to TB)
- Data is **rarely updated in place**, **update through appends**
- Reads and appends are common

# Distributed File System

- Reliable distributed file system
- Data kept in “chunks” spread across multiple machines
- Each chunk replicated on different machines
  - Seamless recovery from disk or machine failure



Chunk servers also serve as compute servers

Bring computation directly to the data!

# Map-Reduce - Restaurant Analogy




# Traditional Scenario


edureka!

Traditional Scenario:

2 orders per hour



Single Cook



Food Shelf

Traditional Scenario:

Data is generated at a steady rate and is structured in nature



Traditional Processing System




RDBMS


# Big Data Scenario

## Scenario 2:

- They started taking Online orders
- 10 orders per hour



Single Cook  
(Regular Computing System)



Food Shelf  
(Data)


edureka!

## Big Data Scenario:

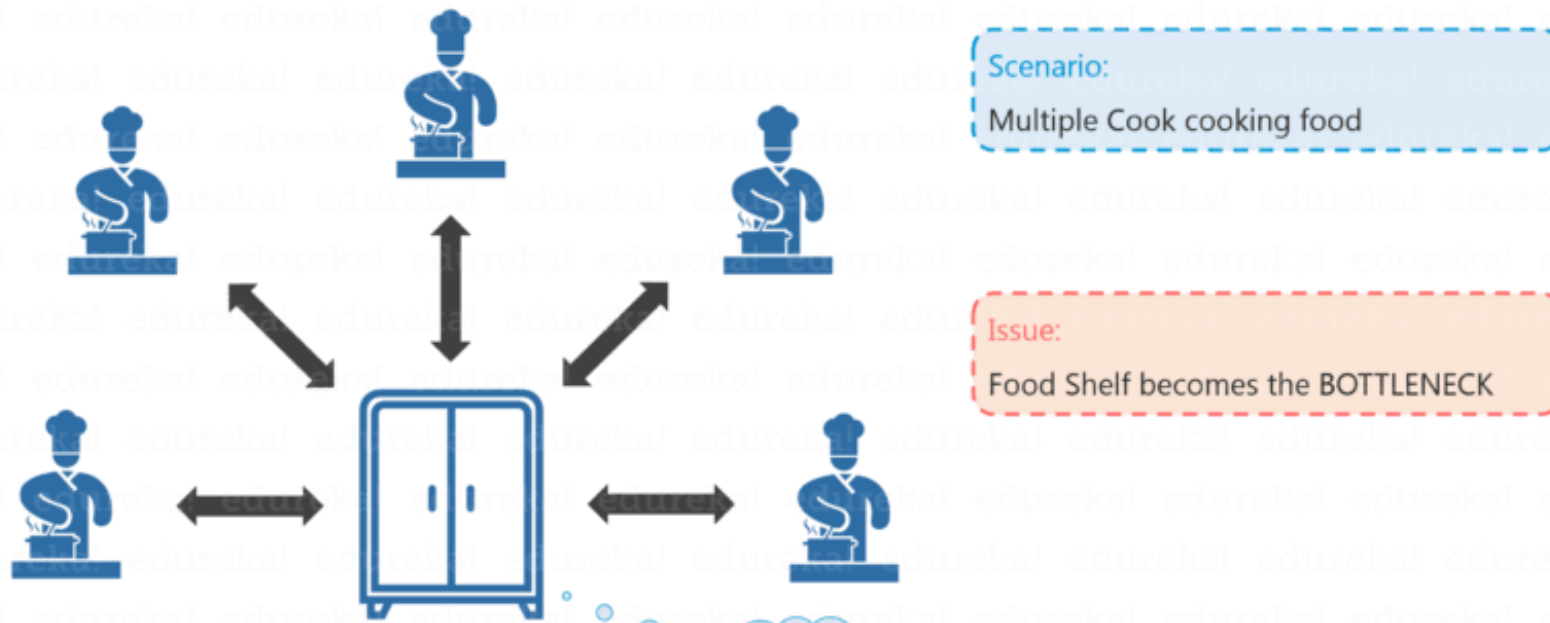
Heterogenous data is being generated at an alarming rate by multiple sources



Traditional Processing  
System



RDBMS



Food Shelf  
(Data)

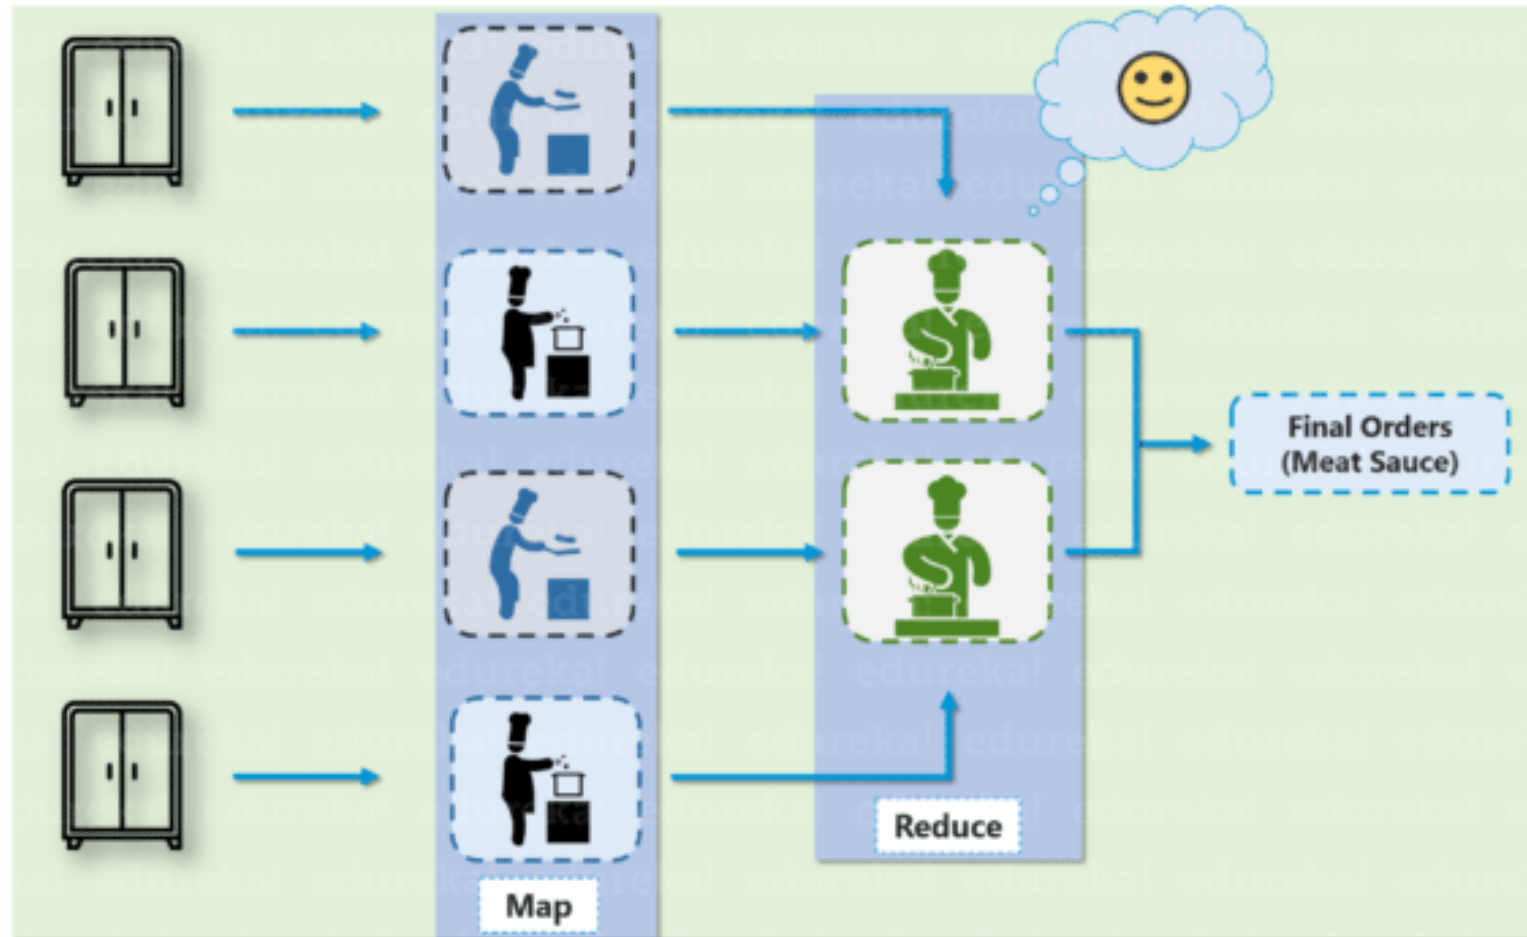


Data Warehouse

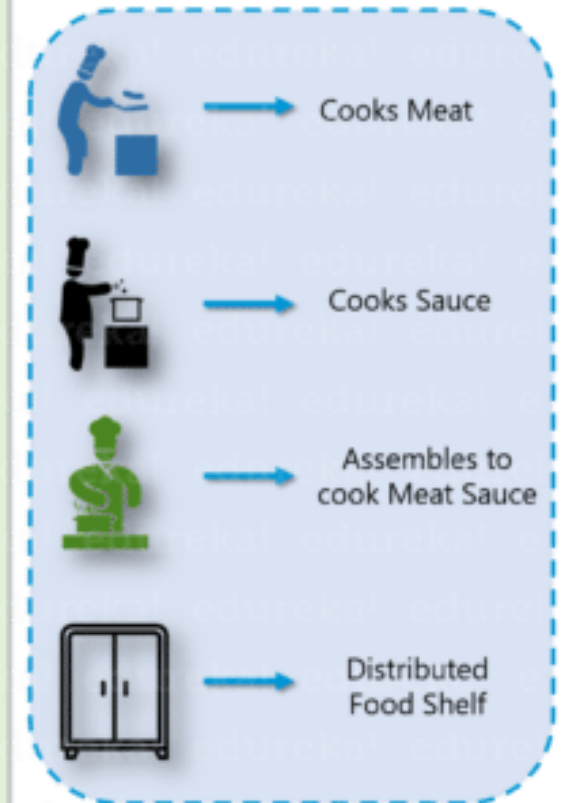
edureka!



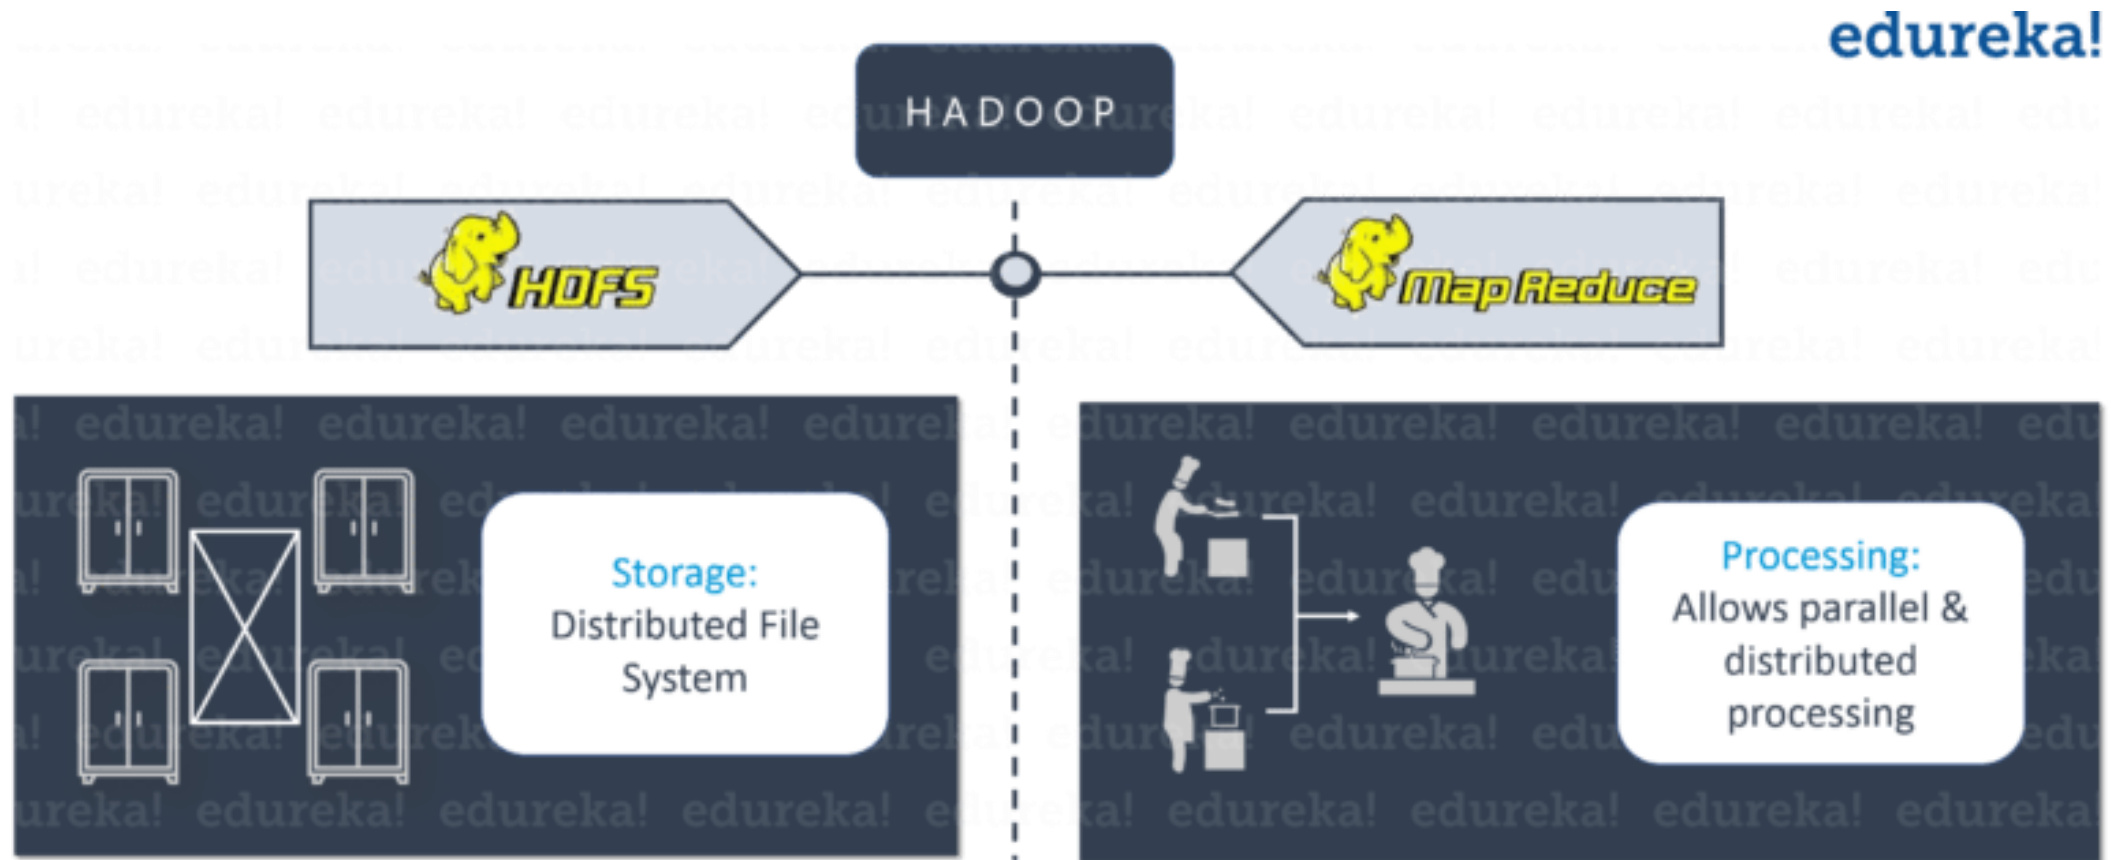
# Apply Map-Reduce Concept



edureka!



# MapReduce and HDFS





## **Map-Reduce Computational Model**

### **Classic Example**

**Count word occurrences in a set of documents**

# Programming Model: MapReduce

- Warm-up task:
  - We have a **huge** text document (e.g., 10 TB)
  - Count the number of times **each distinct word** appears in the file
- Sample application:
  - Analyse web server logs to find popular URLs
  - Term statistics for search engine

# Task: Word Count

## Scenario 1:

- File is too large for RAM, but all `<word, count>` pairs fit in memory

## Scenario 2:

- So many distinct words, and `<word, count>` pairs do not fit in memory
- Divide a large file to small pieces
- Compute in parallel and generate multiple `<word, count>` pairs, and merge them
- Scenario 2 captures the essence of MapReduce

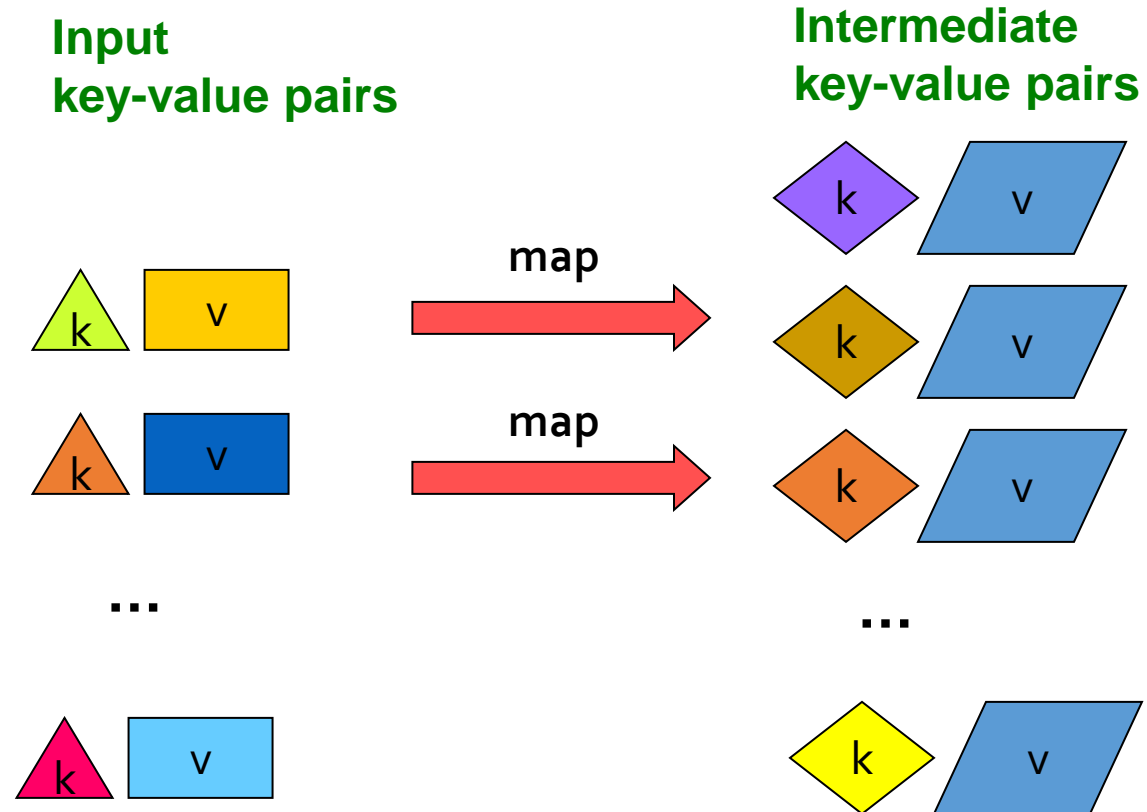
# MapReduce: Overview

- **Map:**
  - Scan input file, e.g., doc.txt
  - Extract something that you care about
- **Group by key:**
  - Sort and Shuffle
- **Reduce:**
  - Aggregate, summarize, filter or transform
  - Write the result


Change Map and Reduce to fit the problem



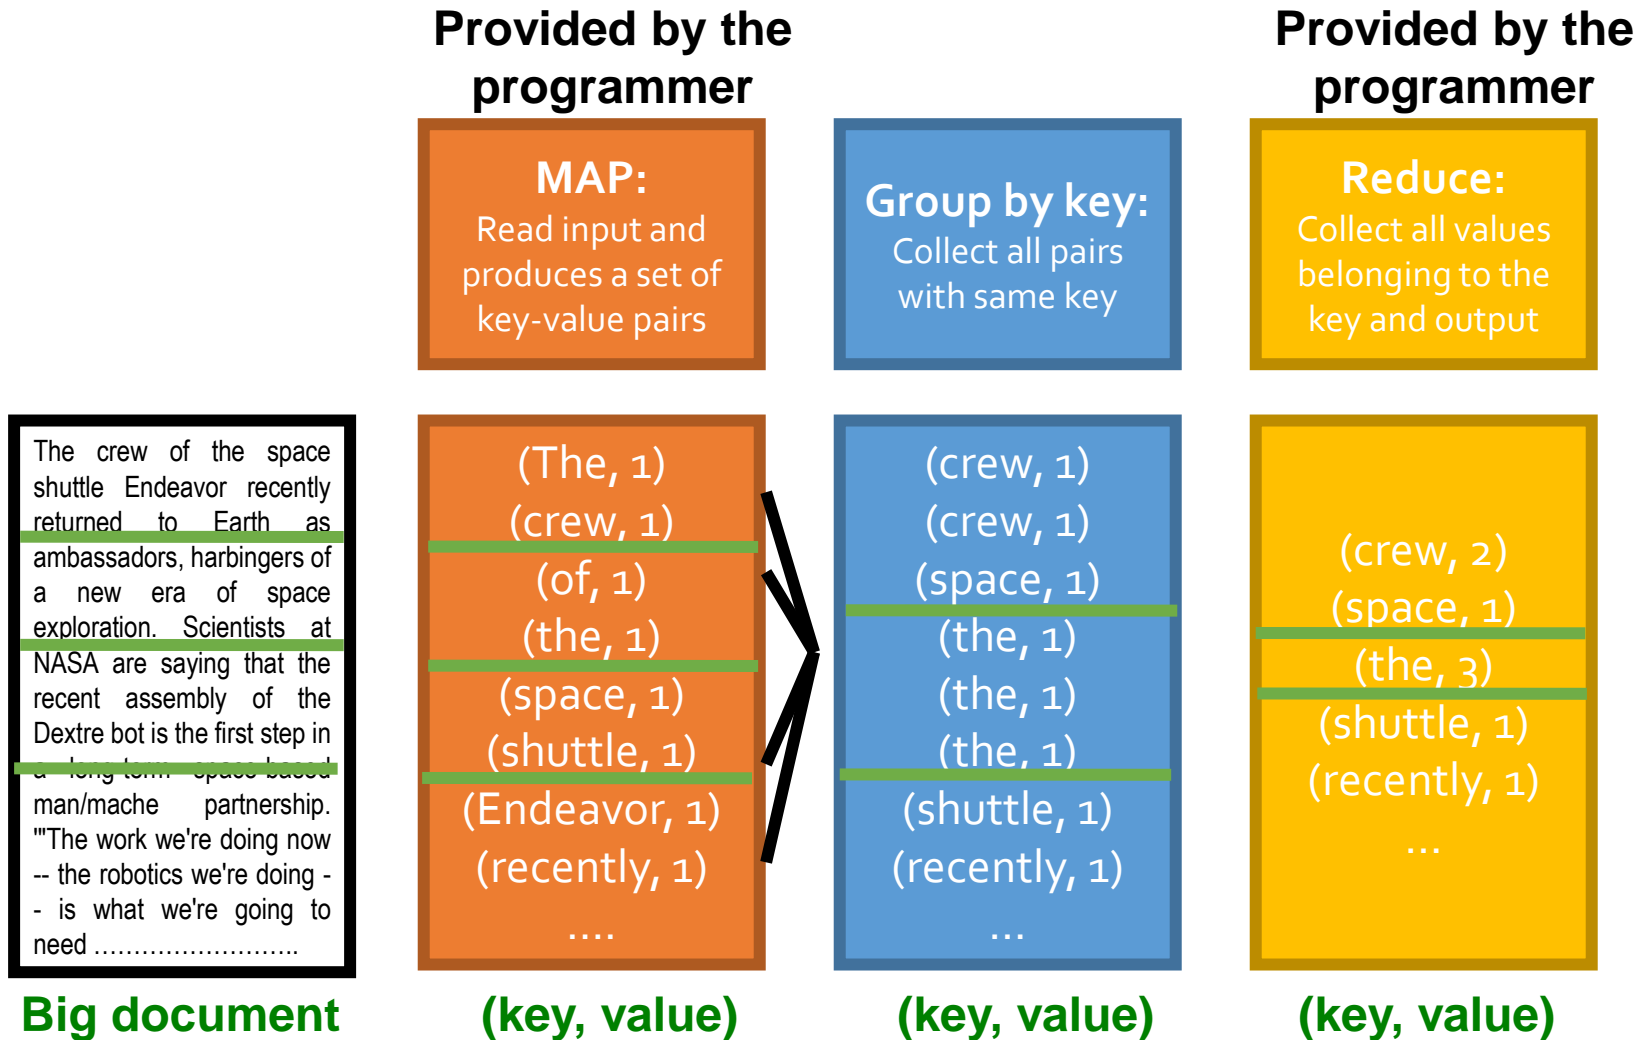
# MapReduce: The Map Step



# MapReduce: The Reduce Step



# MapReduce: Word Counting



# Java Implementation of Word Count MapReduce

```
public class MapClass extends Mapper<LongWritable,
Text, Text, IntWritable> {

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    @Override
    protected void map(LongWritable key, Text value,
Context context) throws IOException,
InterruptedException {

        String line = value.toString();
        StringTokenizer st = new StringTokenizer(line, " ");
        while (st.hasMoreTokens()) {
            word.set(st.nextToken());
            context.write(word, one);
        }
    }
}
```

```
public class ReduceClass extends Reducer<Text,
IntWritable, Text, IntWritable> {

    @Override
    protected void reduce(Text key, Iterable<IntWritable>
values, Context context) throws IOException,
InterruptedException {

        int sum = 0;
        Iterator<IntWritable> valuesIt = values.iterator();
        while (valuesIt.hasNext()) {
            sum = sum + valuesIt.next().get();
        }
        context.write(key, new IntWritable(sum));
    }
}
```

# MapReduce Example: Host Size

- Suppose we have a large web corpus with a metadata file formatted as follows:
  - Each record of the form: (URL, size, date, ...)
- For each host, find out the total number of bytes
- Map
  - For each record, output (hostname (URL), size)
- Reduce
  - Sum the sizes of each host

# MapReduce Example: Language Model

- Count number of times each 5-word sequence occurs in a large corpus of documents
- Map
  - Extract (5-word sequence, count) from document
- Reduce
  - Combine the counts



**Hadoop**



# What is Hadoop

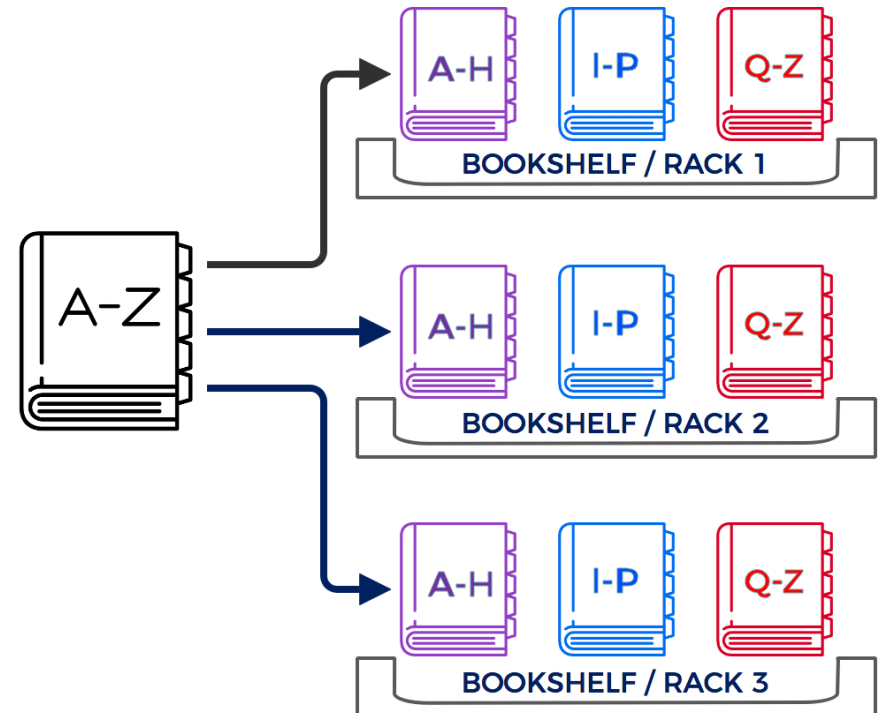
- The Apache Hadoop software library is a **framework** that allows for the **distributed processing of large data** sets across clusters of computers using simple programming models.
  - Hadoop was developed, based on the paper written by Google on the MapReduce system and it applies concepts of functional programming.
  - Hadoop is written in the Java programming language and ranks among the highest-level [Apache projects](#).
- The Hadoop Framework
  - Well-known in big data spaces.
  - Consist of multiple projects of Apache Software Foundation.
  - Support various types of datasets, e.g., structured and unstructured.

# Four Primary Components

- Hadoop Common
  - common utilities that support other modules.
- Hadoop Distributed File System (HDFS)
  - a distributed filesystem that provides high-throughput access to application data.
- Hadoop YARN
  - a framework for job scheduling and cluster resource management.
- Hadoop MapReduce
  - a programming model for parallel processing of large datasets.

# Hadoop Distributed File System (HDFS)

- HDFS performs two main functions
  - **Namespaces**: Provides namespaces that hold cluster metadata, that is, the location of data in the Hadoop cluster
  - **Data storage**: Acts as storage for data used in Hadoop cluster
- Bookshelf Analogy
  - Consider a large book that consists of Chapter A-Z. in HDFS, books have been split into smaller chunks.



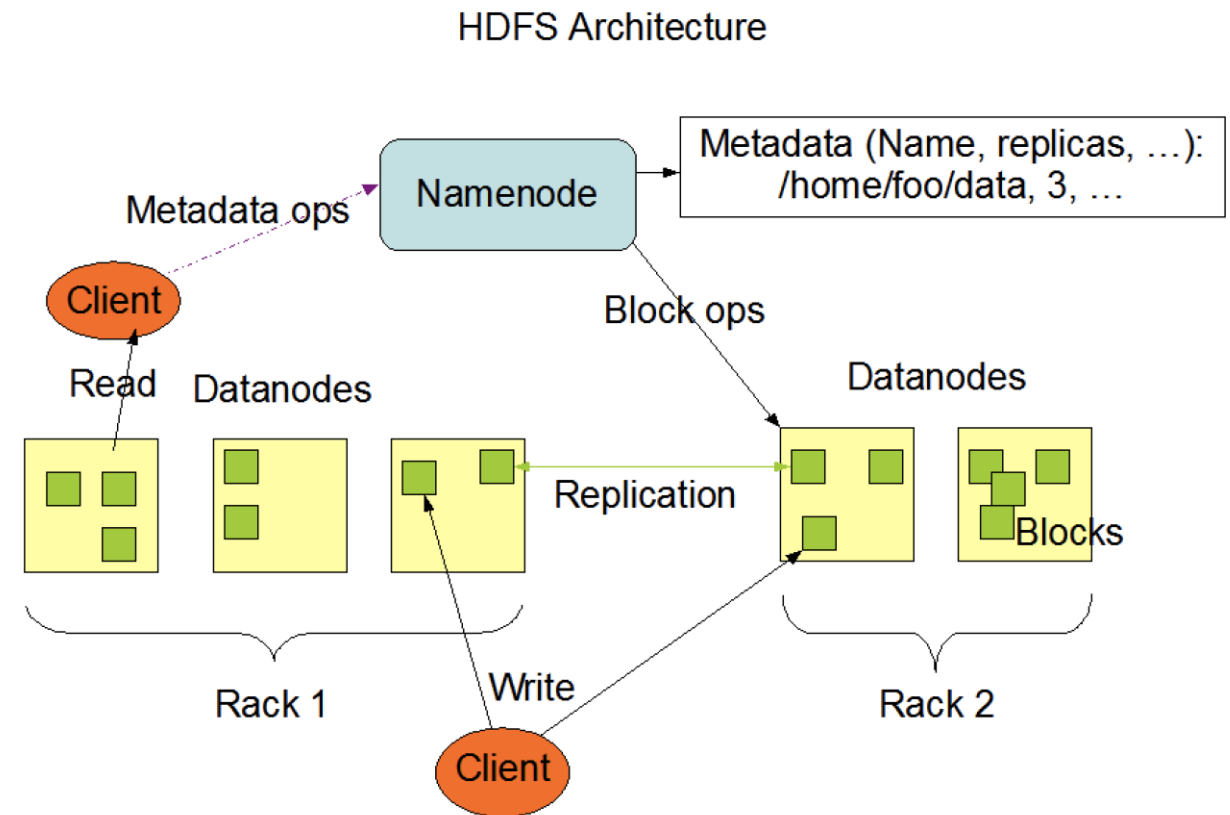
# HDFS Backend

- **NameNode**

- Consider as master node.
- Contain cluster metadata and the location of data
- Store the entire namespace in RAM

- **DataNode**

- Individual servers, responsible for storing chunks of data
- Perform compute operations



# Reference

- Rajaraman, Anand, and Jeffrey David Ullman. *Mining of massive datasets*. Cambridge University Press, 2011.
- Dasgupta, Nataraj. *Practical big data analytics: Hands-on techniques to implement enterprise analytics and machine learning using Hadoop, Spark, NoSQL and R*. Packt Publishing Ltd, 2018.
- Shubham Sinha, *Hadoop Tutorial - A Comprehensive Guide To Hadoop*, Oct 9 2014
- edureka!, <https://www.edureka.co/>
- Leskovec, Jure, Anand Rajaraman, and Jeff Ullman. "Big Data Management and Analytics." (2015).