

# COMP810 Data Warehousing and Big Data

Semester 2 2024

Dr Victor Miranda



# COMP810

## Week 5 Data Warehousing

- Complex OLAP operations  
& SQL queries

# JOINS IN SQL

---

- An SQL 'join' is used fetch data from two or more tables, which is set to appear as a single set of data.
- Right outer join:
  - In practice, a 'join' combines columns from two or more table by using common identifiers (IDs) to both tables
- Keyword: JOIN ... ON

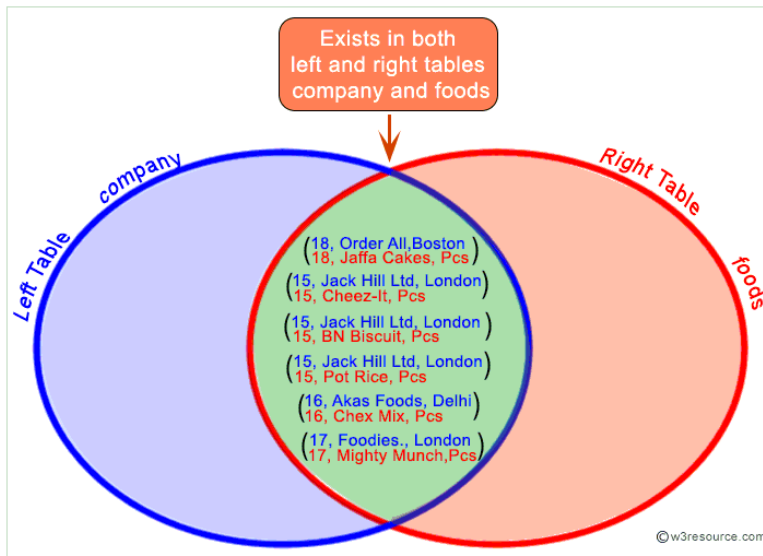
# Types of 'JOIN'

---

- Inner join:
  - Include all the rows from both tables
- Left outer join:
  - Include the left TABLE even if there's no match.
- Right outer join:
  - Include the right TABLE even if there's no match
- Full outer join:
  - Include the both left and right TABLES even if there's no match
- The CROSS join


# The INNER join

- The INNER join selects all COMMON rows from both tables even if there's no match.



# The INNER join

- The INNER join selects all COMMON rows from both tables even if there's no match.



## SQL INNER JOIN

Table: Customers

customer_id	first_name
1	John
2	Robert
<u>3</u>	David
4	John
<u>5</u>	Betty

Table: Orders

order_id	amount	customer
1	200	10
2	500	<u>3</u>
3	300	6
4	800	<u>5</u>
5	150	8

customer_id	first_name	amount
3	David	500
5	Betty	800

# The INNER join - EXAMPLE

A	M
1	m
2	n
4	o

table\_A

A	N
2	p
3	q
5	r


table\_B

```
SELECT * FROM table_A  
INNER JOIN table_B  
ON table_A.A=table_B.A;
```




A	M	A	N
2	n	2	p

Output



# The LEFT join or LEFT OUTER join

- The LEFT JOIN joins two tables and fetches rows based on 'conditions' (matching/linking both tables). There is an 'order': Left table and right table.
- The unmatched rows from the **RIGHT table** will also be available as NULLs if there is no match.





# The LEFT join - EXAMPLE

A	M
1	m
2	n
4	o

table\_A

A	N
2	p
3	q
5	r


table\_B

```
SELECT * FROM table_A  
LEFT JOIN table_B  
ON table_A.A=table_B.A;
```



A	M	A	N
2	n	2	p
1	m	null	null
4	o	null	null

Output



table\_A table\_B


# The LEFT join – EXAMPLE II

Left Table

Date	CountryID	Units
1/1/2020	1	40
1/2/2020	1	25
1/3/2020	3	30
1/4/2020	4	35

Right Table

ID	Country
1	USA
2	Canada
3	Panama




Merged Table

Date	CountryID	Units	Country
1/1/2020	1	40	USA
1/2/2020	1	25	USA
1/3/2020	3	30	Panama
1/4/2020	4	35	<i>null</i>

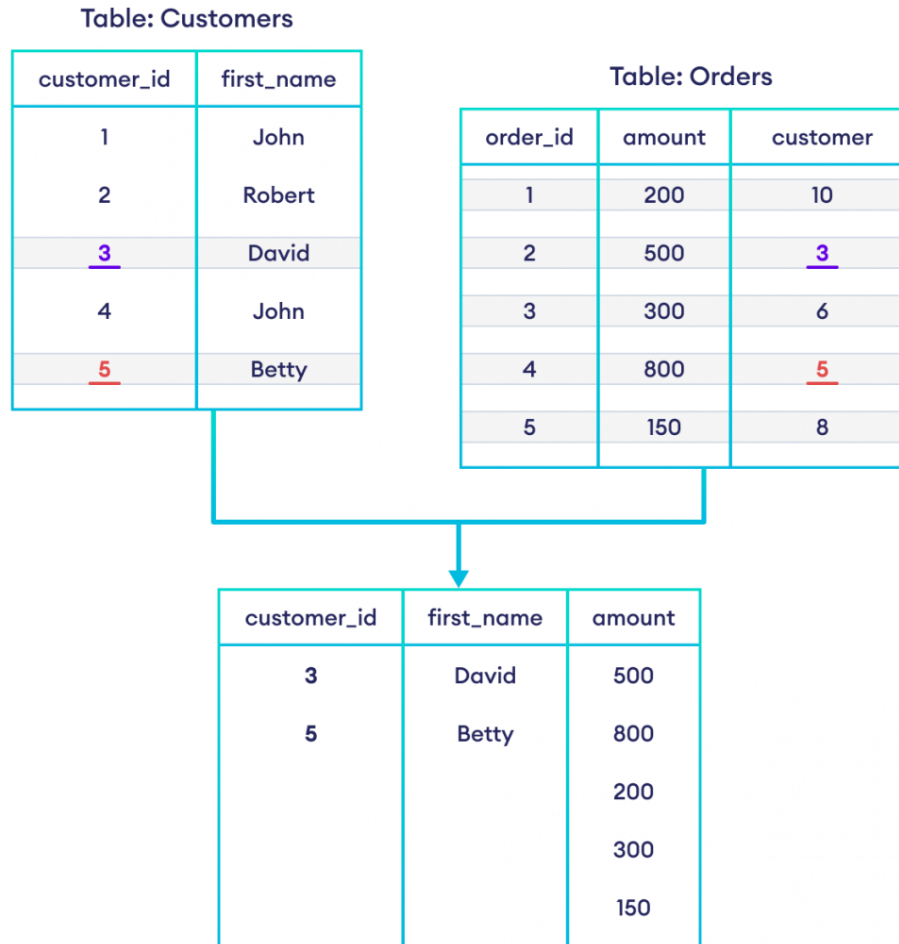
# The RIGHT join

- The RIGHT JOIN joins two tables and fetches rows based on 'conditions' (matching/linking both tables). There is an 'order': Left table and right table.
- The unmatched tables from the **LEFT table** will also be available as NULLs if there is no match.



# The RIGHT join – EXAMPLE

## SQL RIGHT JOIN



# The RIGHT join – EXAMPLE II

```
SELECT * FROM table_A  
RIGHT JOIN table_B  
ON table_A.A=table_B.A;
```



A	M
1	m
2	n
4	o


table\_A

A	N
2	p
3	q
5	r

table\_B

A	M	A	N
2	n	2	p
null	null	3	q
null	null	5	r

Output




table\_A

table\_B

# The FULL join

- The FULL JOIN returns all records when there is a match in left (Table 1) **or** right (Table 2) records.



# The FULL join

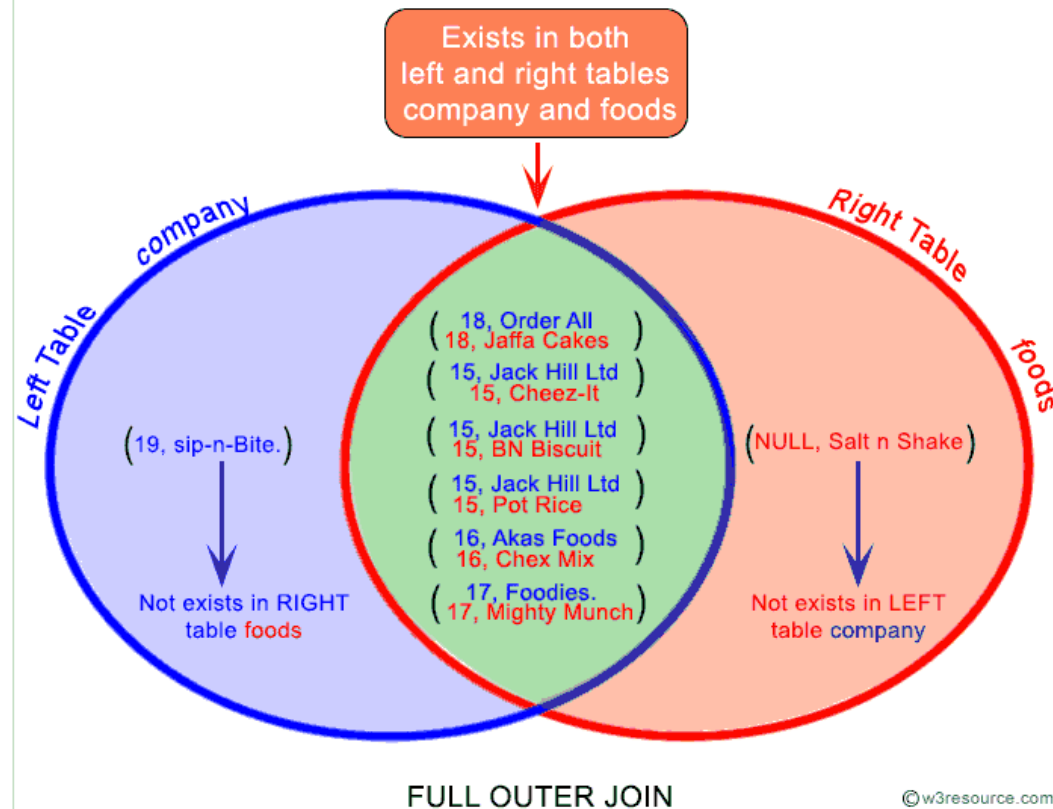
- The FULL JOIN returns all records when there is a match in left (Table 1) **or** right (Table 2) records.

```
SELECT a.company_id as "a.ComID",
a.company_name as "C_Name",
b.company_id as "b.ComID"
FROM company a
LEFT OUTER JOIN foods b
ON a.company_id = b.company_id;
```

A.ComID	C_Name	B.ComID
16	Akas Foods	16
15	Jack Hill Ltd	15
15	Jack Hill Ltd	15
17	Foodies.	17
15	Jack Hill Ltd	15
18	Order All	18
19	sip-n-Bite.	-

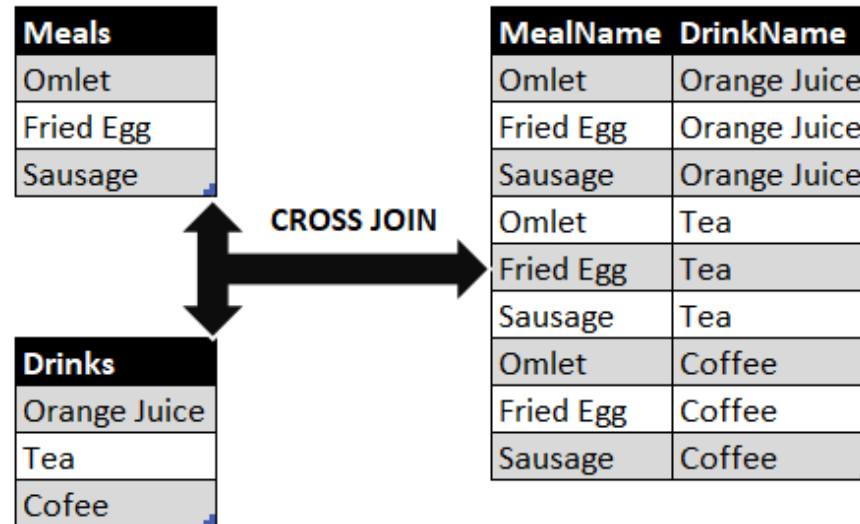
```
SELECT a.company_id as "a.ComID",
a.company_name as "C_name",
b.company_id as "b.ComID"
FROM company a
RIGHT OUTER JOIN foods b
ON a.company_id = b.company_id;
```

A.ComID	C_name	B.ComID
16	Akas Foods	16
15	Jack Hill Ltd	15
15	Jack Hill Ltd	15
17	Foodies.	17
15	Jack Hill Ltd	15
18	Order All	18
-	-	-



# The CROSS join

- The CROSS JOIN produces crossed-set, where all entries from the first table (Table 1) are 'crossed' with all entries from the second table (also called Cartesian product).
- If there is a WHERE clause, this call resembles an INNER JOIN





# The CROSS join - EXAMPLE


A	M
1	m
2	n
4	o

table\_A

A	N
2	p
3	q
5	r

table\_B

**SELECT \***  
**FROM table\_A**  
**CROSS JOIN table\_B;**



A	M	A	N
1	m	2	p
2	n	2	p
4	o	2	p
1	m	3	q
2	n	3	q
4	o	3	q
1	m	5	r
2	n	5	r
4	o	5	r

Output

---

## More examples

# SQL– Left/Right Join

## Example:

### Schema:

People(person\_fname, person\_lname, person\_id, person\_state, person\_city)

Movies(movie\_id, movie\_title, director\_id, studio\_id)

Location(movie\_id, city, state)

### Query:

Select movie\_title, city, state

From Movies Left Join Locations

On Movies.movie\_id = Locations.movie\_id

Select movie\_title, person\_fname, person\_lname

From Movies Right Join People

On Movies.director\_id = Person.person\_id

# SQL– Left/Right Join

## Example:

### Schema:

People(person\_fname, person\_lname, person\_id, person\_state, person\_city)

Movies(movie\_id, movie\_title, director\_id, studio\_id)

Location(movie\_id, city, state)

### Query:

Select movie\_title, city, state

From Movies Left Join Locations

On Movies.movie\_id = Locations.movie\_id

**Includes all  
matched  
movie titles**

Select movie\_title, person\_fname, person\_lname

From Movies Right Join People

On Movies.director\_id = Person.person\_id

**Includes  
all people  
matching  
to directors**

# SQL Inner Joins

```
SELECT S.name, E.classid  
FROM students S (INNER) JOIN Enrolled E  
ON S.sid=E.sid
```

**S**

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Jones	DataScience194
Smith	French150

**E**

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10

Note the previous version of this query (with no join keyword) is an “Implicit join”

# SQL Inner Joins

```
SELECT S.name, E.classid  
FROM students S (INNER) JOIN Enrolled E  
ON S.sid=E.sid
```

**S**

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Jones	DataScience194
Smith	French150

**E**

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10

Unmatched keys

# What kind of Join is this?

```
SELECT S.name, E.classid
FROM Students S ?? Enrolled E
ON S.sid=E.sid
```

**S**

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Jones	DataScience194
Smith	French150
Brown	NULL

**E**

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10

# SQL Joins

```
SELECT S.name, E.classid
FROM students S LEFT OUTER JOIN Enrolled E
ON S.sid=E.sid
```

**S**

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Jones	DataScience194
Smith	French150
Brown	NULL

**E**

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10



# What kind of Join is this?

```
SELECT S.name, E.classid
FROM students S ?? Enrolled E
ON S.sid=E.sid
```

**S**

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Jones	DataScience194
Smith	French150
NULL	English10

**E**

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10

# SQL Joins

```
SELECT S.name, E.classid
FROM Students S RIGHT OUTER JOIN Enrolled E
ON S.sid=E.sid
```

**S**

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Jones	DataScience194
Smith	French150
NULL	English10

**E**

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10

# SQL Joins

```
SELECT S.name, E.classid
FROM students S ? JOIN Enrolled E
ON S.sid=E.sid
```

**S**

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Jones	DataScience194
Smith	French150
NULL	English10
Brown	NULL

**E**

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10

# SQL Joins

```
SELECT S.name, E.classid  
FROM Students S FULL OUTER JOIN Enrolled E  
ON S.sid=E.sid
```

**S**

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Jones	DataScience194
Smith	French150
NULL	English10
Brown	NULL

**E**

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10

# What kind of Join is this?

```
SELECT S.name, E.classid
FROM Students S ?? Enrolled E
ON S.sid=E.sid
```

**S**

S.name	S.sid
Jones	11111
Smith	22222
Brown	33333

S.name	E.classid
Jones	History105
Smith	French150

**E**

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150
44444	English10

# SQL Joins

```
SELECT S.name, E.classid
FROM Students S LEFT SEMI JOIN Enrolled E
ON S.sid=E.sid
```

<b>S</b>	S.name	S.sid
	Jones	11111
	Smith	22222
	Brown	33333

S.name	E.classid
Jones	History105
Smith	French150

<b>E</b>	E.sid	E.classid
	11111	History105
	11111	DataScience194
	22222	French150
	44444	English10

# What kind of Join is this?

SELECT \*

FROM students S ?? Enrolled E

**S**

S.name	S.sid
Jones	11111
Smith	22222

**E**

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150

S.name	S.sid	E.sid	E.classid
Jones	11111	11111	History105
Jones	11111	11111	DataScience194
Jones	11111	22222	French150
Smith	22222	11111	History105
Smith	22222	11111	DataScience194
Smith	22222	22222	French150

# SQL Joins

SELECT \*

FROM **students S CROSS JOIN Enrolled E**

**S**

S.name	S.sid
Jones	11111
Smith	22222

**E**

E.sid	E.classid
11111	History105
11111	DataScience194
22222	French150


S.name	S.sid	E.sid	E.classid
Jones	11111	11111	History105
Jones	11111	11111	DataScience194
Jones	11111	22222	French150
Smith	22222	11111	History105
Smith	22222	11111	DataScience194
Smith	22222	22222	French150




---

# Summary


# SQL JOINS




```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
```




```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
```




```
SELECT <select_list>
FROM TableA A
INNER JOIN TableB B
ON A.Key = B.Key
```




```
SELECT <select_list>
FROM TableA A
LEFT JOIN TableB B
ON A.Key = B.Key
WHERE B.Key IS NULL
```



```
SELECT <select_list>
FROM TableA A
RIGHT JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
```



```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
```



```
SELECT <select_list>
FROM TableA A
FULL OUTER JOIN TableB B
ON A.Key = B.Key
WHERE A.Key IS NULL
OR B.Key IS NULL
```

---

# The UNION operator

# The UNION operator


---

It is used to 'combine' the resulting set of two or more SELECT operations

## NOTES:

- (a) Every SELECT output must have the same number of columns
- (b) The columns should have similar data types
- (c) The columns from each SELECT operation must be sorted equally.

# The UNION operator - EXAMPLE



---

# SQL Views

# Types of Views

---

## ■ Virtual views:


- Used in databases
- Computed only on-demand – *slower* at runtime
- Always up to date

## ■ Materialized views

- Used in data warehouses
- Precomputed offline – *faster* at runtime
- May have stale data

# SQL Views

- In SQL, a view is a virtual table based on the result-set of an SQL statement. A view contains rows and columns, just like a real table.





# Defining Views - Syntax

Views are relations, except that they are not physically stored.

## CREATE VIEW Syntax

```
CREATE VIEW view_name AS  
SELECT column1, column2, ...  
FROM table_name  
WHERE condition;
```


# Defining Views - Syntax

Views are relations, except that they are not physically stored.

For presenting different information to different users

**Employee**(ssn, name, department, project, salary)

```
CREATE VIEW Developers AS
  SELECT name, project
  FROM Employee
  WHERE department = "Development"
```



Payroll has access to **Employee**, others only to **Developers**

# A Different View

Person(name, city)

Purchase(buyer, seller, product, store)

Product(name, maker, category)

```
CREATE VIEW Seattle-view AS

SELECT buyer, seller, product, store
FROM Person, Purchase
WHERE Person.city = "Seattle" AND
      Person.name = Purchase.buyer
```

We have a new virtual table:

Seattle-view(buyer, seller, product, store)

# A Different View

We can later use the view:

```
SELECT name, store  
FROM Seattle-view, Product  
WHERE Seattle-view.product = Product.name AND  
Product.category = "shoes"
```

# What Happens When We Query a View ?

```
SELECT name, Seattle-view.store
FROM   Seattle-view, Product
WHERE  Seattle-view.product = Product.name AND
       Product.category = "shoes"
```



```
SELECT name, Purchase.store
FROM   Person, Purchase, Product
WHERE  Person.city = "Seattle" AND
       Person.name = Purchase.buyer AND
       Purchase.poduct = Product.name AND
       Product.category = "shoes"
```

# Updating Views

How can I insert a tuple into a table that doesn't exist?

**Employee**(ssn, name, department, project, salary)

```
CREATE VIEW Developers AS  
  SELECT name, project  
  FROM Employee  
  WHERE department = "Development"
```

If we make the  
following insertion:

```
INSERT INTO Developers  
VALUES("Joe", "Optimizer")
```

It becomes:

```
INSERT INTO Employee  
VALUES(NULL, "Joe", NULL, "Optimizer", NULL)
```

# Non-Updatable Views

```
CREATE VIEW Seattle-view AS  
  
    SELECT seller, product, store  
    FROM   Person, Purchase  
    WHERE  Person.city = "Seattle" AND  
           Person.name = Purchase.buyer
```

How can we add the following tuple to the view?

("Joe", "Shoe Model 12345", "Nine West")

We need to add "Joe" to Person first, but we don't have all its attributes

# Week 6 ( Week 12)

Materialized Views

More Complex OLAP operations with SQL



# References:

---

(a) A Conceptual Poverty Mapping Data Model

Link: [https://www.researchgate.net/figure/Key-thematic-layers-for-poverty-spatial-data-modeling\\_fig2\\_229724703](https://www.researchgate.net/figure/Key-thematic-layers-for-poverty-spatial-data-modeling_fig2_229724703)

(b) Relational Database relationships

<https://www.youtube.com/watch?v=C3icLzBtg8I>

(c) <https://courses.ischool.berkeley.edu/i202/f97/Lecture13/DatabaseDesign/sld002.htm>

(d) <https://nexwebsites.com/database/database-management-systems/>

(e) Acknowledgement – Thanks to <http://courses.cs.washington.edu/courses/cse544/> for providing part of this presentation.

(f) Acknowledgement – Thanks to © Silberchatz, Korth and Surdashaan for providing part of this presentation.

(e) Malinowski, Elzbieta, Zimányi, Esteban (2008) *Advanced Data Warehouse Design: From Conventional to Spatial and Temporal Applications*. Springer Berlin Heidelberg. Copyright © 2008 Elzbieta Malinowski & Esteban Zimányi